



Sardar Patel Institute of Technology, Mumbai  
Department of Electronics and Telecommunication Engineering  
T.E. Sem-V (2018-2019)

ETL54-Statistical Computational Laboratory

**Lab-4: Classification Analysis and Modeling**

**Name: Manish Dsilva**

**Roll No.15**

**Objective:** To carry out logistic regression (including multiple regression and multiclass regression) and build a regression model

**Outcomes:**

1. Part-I: Implement Logistic Regression for University Admission dataset :  
Target variable: Admit or reject  
Predictors or independent variables: GRE, RANK, GPA
2. Part-II: Implement Multinomial Logistics Regression for iris

**System Requirements:** Ubuntu OS with R and RStudio installed

**Logistic Regression**

We use the logistic regression equation to predict the probability of a dependent variable taking the dichotomy values 0 or 1. Suppose  $x_1, x_2, \dots, x_p$  are the independent variables,  $\alpha$  and  $\beta_k$  ( $k = 1, 2, \dots, p$ ) are the parameters, and  $E(y)$  is the expected value of the dependent variable  $y$ , then the logistic regression equation is:

$$E(y) = 1 / (1 + e^{-(\alpha + \sum_k \beta_k x_k)})$$

For example, in the built-in data set *mtcars*, the data column *am* represents the transmission type of the automobile model (0 = automatic, 1 = manual).

With the logistic regression equation, we can model the probability of a manual transmission in a vehicle based on its engine horsepower and weight data.

$$P(\text{Manual Transmission}) = 1 / (1 + e^{-(\alpha + \beta_1 * \text{Horsepower} + \beta_2 * \text{Weight})})$$

**Estimated Logistic Regression Equation**

Using the generalized linear model, an estimated logistic regression equation can be formulated as below. The coefficients  $a$  and  $b_k$  ( $k = 1, 2, \dots, p$ ) are determined according to a maximum likelihood approach, and it allows us to estimate the probability of the dependent variable  $y$  taking on the value 1 for given values of  $x_k$  ( $k = 1, 2, \dots, p$ ).

$$\text{Estimate of } P(y = 1 \mid x_1, \dots, x_p) = 1 / (1 + e^{-(a + \sum_k b_k x_k)})$$

We apply the function **glm** to a formula that describes the transmission type (am) by the horsepower (hp) and weight (wt). This creates a generalized linear model (GLM) in the binomial family.

In R:

#Build a model:

```
am.glm = glm(formula=am ~ hp + wt, data=mtcars, family=binomial)
```

#Test data

```
newdata = data.frame(hp=120, wt=2.8)
```

#Predict

```
predict(am.glm, newdata, type="response")
```

Let's start with an **example confusion matrix for a binary classifier** (though it can easily be extended to the case of more than two classes):

n=165		Predicted: NO	Predicted: YES
Actual: NO		50	10
Actual: YES		5	100

What can we learn from this matrix?

- There are two possible predicted classes: "yes" and "no". If we were predicting the presence of a disease, for example, "yes" would mean they have the disease, and "no" would mean they don't have the disease.
- The classifier made a total of 165 predictions (e.g., 165 patients were being tested for the presence of that disease).
- Out of those 165 cases, the classifier predicted "yes" 110 times, and "no" 55 times.
- In reality, 105 patients in the sample have the disease, and 60 patients do not.

Let's now define the most basic terms, which are whole numbers (not rates):

- **true positives (TP):** These are cases in which we predicted yes (they have the disease), and they do have the disease.
- **true negatives (TN):** We predicted no, and they don't have the disease.

- **false positives (FP):** We predicted yes, but they don't actually have the disease. (Also known as a "Type I error.")
- **false negatives (FN):** We predicted no, but they actually do have the disease. (Also known as a "Type II error.")

I've added these terms to the confusion matrix, and also added the row and column totals:

		Predicted:		
		NO	YES	
Actual:	NO	TN = 50	FP = 10	60
	YES	FN = 5	TP = 100	105
		55	110	

This is a list of rates that are often computed from a confusion matrix for a binary classifier:

- **Accuracy:** Overall, how often is the classifier correct?
  - $(TP+TN)/total = (100+50)/165 = 0.91$
- **Misclassification Rate:** Overall, how often is it wrong?
  - $(FP+FN)/total = (10+5)/165 = 0.09$
  - equivalent to 1 minus Accuracy
  - also known as "Error Rate"
- **True Positive Rate:** When it's actually yes, how often does it predict yes?
  - $TP/actual\ yes = 100/105 = 0.95$
  - also known as "Sensitivity" or "Recall"
- **False Positive Rate:** When it's actually no, how often does it predict yes?
  - $FP/actual\ no = 10/60 = 0.17$
- **True Negative Rate:** When it's actually no, how often does it predict no?
  - $TN/actual\ no = 50/60 = 0.83$
  - equivalent to 1 minus False Positive Rate
  - also known as "Specificity"
- **Precision:** When it predicts yes, how often is it correct?
  - $TP/predicted\ yes = 100/110 = 0.91$
- **Prevalence:** How often does the yes condition actually occur in our sample?
  - $actual\ yes/total = 105/165 = 0.64$

A couple other terms are also worth mentioning:

- **Null Error Rate:** This is how often you would be wrong if you always predicted the majority class. (In our example, the null error rate would be  $60/165=0.36$  because if you always predicted yes, you would only be wrong for the 60 "no" cases.) This can be a useful baseline metric to compare your classifier against. However, the best classifier for a particular application will sometimes have a higher error rate than the null error rate, as demonstrated by the [Accuracy Paradox](#)
- **F Score:** This is a weighted average of the true positive rate (recall) and precision.

## Output: Binary Logistic Regression

### Step 1: Data Load and Pre-Processing

```
> mydata <- read.csv("~/Downloads/binary.csv",header=T)
```

```
> str(mydata)
```

```
'data.frame': 400 obs. of 4 variables:
```

```
$ admit: int 0 1 1 1 0 1 1 0 1 0 ...
```

```
$ gre : int 380 660 800 640 520 760 560 400 540 700 ...
```

```
$ gpa : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
```

```
$ rank : int 3 3 1 4 4 2 1 2 3 2 ...
```

```
> mydata$admit <- as.factor(mydata$admit)
```

```
> mydata$rank <- as.factor(mydata$rank)
```

```
> str(mydata)
```

```
'data.frame': 400 obs. of 4 variables:
```

```
$ admit: Factor w/ 2 levels "0","1": 1 2 2 2 1 2 2 1 2 1 ...
```

```
$ gre : int 380 660 800 640 520 760 560 400 540 700 ...
```

```
$ gpa : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
```

```
$ rank : Factor w/ 4 levels "1","2","3","4": 3 3 1 4 4 2 1 2 3 2 ...
```

```
> xtabs(~admit+rank,data=mydata)
```

```
rank
```

```
admit 1 2 3 4
```

```
0 28 97 93 55
```

```
1 33 54 28 12
```

### Step 2: Split Train and Test Data

```

> set.seed(1234)
> ind <- sample(2,nrow(mydata),replace = T,prob=c(0.8,0.2))
> trainData <- mydata[ind == 1,]
> testData <- mydata[ind == 2,]
> testData
      admit gre  gpa rank
5      0 520 2.93   4
14     0 700 3.08   2
16     0 480 3.44   3
26     1 800 3.66   1
28     1 520 3.74   4

> str(testData)
'data.frame': 75 obs. of 4 variables:
 $ admit: Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 2 1 2 ...
 $ gre  : int  520 700 480 800 520 780 500 520 600 620 ...
 $ gpa  : num  2.93 3.08 3.44 3.66 3.74 3.22 3.13 2.68 2.82 3.18 ...
 $ rank : Factor w/ 4 levels "1","2","3","4": 4 2 3 1 4 2 2 3 4 2 ...

> str(trainData)
'data.frame': 325 obs. of 4 variables:
 $ admit: Factor w/ 2 levels "0","1": 1 2 2 2 2 2 1 2 1 1 ...
 $ gre  : int  380 660 800 640 760 560 400 540 700 800 ...
 $ gpa  : num  3.61 3.67 4 3.19 3 2.98 3.08 3.39 3.92 4 ...
 $ rank : Factor w/ 4 levels "1","2","3","4": 3 3 1 4 2 1 2 3 2 4 ...

```

### Step 3: Build a Classifying Model

```

> model <- glm(formula = admit ~ gre + gpa + rank, data = trainData, family = 'binomial' )
> model

```

```

Call: glm(formula = admit ~ gre + gpa + rank, family = "binomial",
  data = trainData)

```

Coefficients:

```

(Intercept)      gre      gpa    rank2    rank3    rank4
-5.009514    0.001631    1.166408  -0.570976  -1.125341  -1.532942

```

Degrees of Freedom: 324 Total (i.e. Null); 319 Residual

Null Deviance: 404.4

Residual Deviance: 370 AIC: 382

```
> summary(model)
```

```
(Intercept) -5.009514  1.316514 -3.805 0.000142 ***
```

```
gre      0.001631  0.001217  1.340 0.180180
```

```
gpa      1.166408  0.388899  2.999 0.002706 **
```

```
rank2    -0.570976  0.358273 -1.594 0.111005
```

```
rank3    -1.125341  0.383372 -2.935 0.003331 **
```

```
rank4    -1.532942  0.477377 -3.211 0.001322 **
```

```
--- Call:
```

```
>glm(formula = admit ~ gre + gpa + rank, family = "binomial",  
      data = trainData)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.5873	-0.8679	-0.6181	1.1301	2.1178

Coefficients:

Estimate	Std. Error	z value	Pr(> z )
----------	------------	---------	----------

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 404.39 on 324 degrees of freedom

Residual deviance: 369.99 on 319 degrees of freedom

AIC: 381.99

Number of Fisher Scoring iterations: 4

Step 4: Rebuilding the model with relevant features

```
> model <- glm(formula = admit ~ gpa + rank,data = trainData,family = 'binomial' )
```

```
> summary(model)
```

```
Call:
```

```
glm(formula = admit ~ gpa + rank, family = "binomial", data = trainData)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.5156	-0.8880	-0.6318	1.1091	2.1688

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-4.7270	1.2918	-3.659	0.000253 ***
gpa	1.3735	0.3590	3.826	0.000130 ***
rank2	-0.5712	0.3564	-1.603	0.108976
rank3	-1.1645	0.3804	-3.061	0.002203 **
rank4	-1.5642	0.4756	-3.289	0.001005 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 404.39 on 324 degrees of freedom

Residual deviance: 371.81 on 320 degrees of freedom

AIC: 381.81

Number of Fisher Scoring iterations: 4

## Step 5: Predicting the model based on Test Data

```
> pred <- predict(model, testData, type = "response")
```

```
> head(pred)
```

5	14	16	26	28	29
0.09390783	0.25582628	0.23746963	0.57446309	0.23971008	0.29411490

```
> pred[pred>0.5]
```

26	90	92	140	158	369	396
0.5744631	0.5488115	0.5677344	0.5474076	0.5474076	0.6828897	0.5488115

```
> p1 <- ifelse(pred>0.5,1,0)
```

```
> head(p1)
```

5	14	16	26	28	29
0	0	0	1	0	0

```
> table(p1)
```

p1

```

0 1
68 7
> table(testData$admit)
0 1
50 25

```

### Step 6: Confusion Matrix

```

> tab <- table(p1,testData$admit)
> tab
p1 0 1
0 48 20
1 2 5
> tab <- table(Predicted =p1,Actual =testData$admit)
> tab
      Actual
Predicted 0 1
0 48 20
1 2 5

```

### Step 7: Model Performance Parameters

Accuracy	0.71	Prevalance	0.3333
Misclass. Rate	0.29	Precision	0.7143
True Pos Rate	0.2	Null Rate Error	0.3333
False Pos Rate	0.04	F-Score	0.3125

## Output: Multiclass Logistic Regression

### Step 1: Data Load and Pre-Processing

```

> mydata <- iris
> str(mydata)
'data.frame': 150 obs. of 5 variables:
 $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...

```



```

$ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 ...
> ind <- sample(2,nrow(mydata),replace = T,prob=c(0.8,0.2))
> str(trainData)

'data.frame': 118 obs. of 5 variables:
 $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
> str(testData)

'data.frame': 32 obs. of 5 variables:
 $ Sepal.Length: num 4.8 4.8 5.8 5.1 4.8 5.5 4.9 5 4.5 5.1 ...
 $ Sepal.Width : num 3.4 3 4 3.7 3.4 4.2 3.6 3.5 2.3 3.8 ...
 $ Petal.Length: num 1.6 1.4 1.2 1.5 1.9 1.4 1.4 1.3 1.3 1.9 ...
 $ Petal.Width : num 0.2 0.1 0.2 0.4 0.2 0.2 0.1 0.3 0.3 0.4 ...
 $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

```

## Step 2: Load Library and Create Model

```

> library(nnet)
> model<-multinom(formula=Species~Sepal.Length+Sepal.Width+Petal.Length+Petal.Width,data=trainData)

# weights: 18 (10 variable)
initial value 129.636250
iter 10 value 10.683012
iter 20 value 5.933903
iter 30 value 5.873500
iter 40 value 5.866866
iter 50 value 5.861992
iter 60 value 5.860395
iter 70 value 5.859634
iter 80 value 5.859340
iter 90 value 5.859208
iter 100 value 5.859118
final value 5.859118
stopped after 100 iterations
> model

Call:

```

```
multinom(formula = Species ~ Sepal.Length + Sepal.Width + Petal.Length +
  Petal.Width, data = trainData)
```

Coefficients:

	(Intercept)	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
versicolor	20.95306	-2.019734	-12.17769	10.47244	-2.626553
virginica	-18.78599	-4.541125	-18.31369	19.53561	14.264642

Residual Deviance: 11.71824

AIC: 31.71824

> summary(model)

Call:

```
multinom(formula = Species ~ Sepal.Length + Sepal.Width + Petal.Length +
  Petal.Width, data = trainData)
```

Coefficients:

	(Intercept)	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
versicolor	20.95306	-2.019734	-12.17769	10.47244	-2.626553
virginica	-18.78599	-4.541125	-18.31369	19.53561	14.264642

Std. Errors:

	(Intercept)	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
versicolor	41.61007	134.8689	191.1050	76.27252	17.95107
virginica	42.22799	134.8825	191.1707	76.48059	19.08432

Residual Deviance: 11.71824

AIC: 31.71824

### Step 3: Predict Based on Model

> pred <- predict(model, testData, type = "class")

> pred

```
[1] setosa setosa setosa setosa setosa setosa setosa setosa
[9] setosa setosa setosa versicolor versicolor versicolor versicolor versicolor
[17] versicolor virginica virginica virginica virginica virginica virginica virginica
[25] virginica virginica virginica virginica virginica virginica virginica virginica
Levels: setosa versicolor virginica
```

#### Step 4: Confusion Matrix

```
> tab <- table(pred,testData$Species)
```

```
> tab
```

```
      pred      setosa versicolor virginica
setosa    11         0          0
versicolor 0         6          0
virginica  0         0         15
```

#### Step 5: Model Test Parameters

Accuracy	1	Prevalance	1
Misclass. Rate	0	Precision	1
True Pos Rate	1	Null Rate Error	0.53
False Pos Rate	0	F-Score	1

#### References:

<https://www.youtube.com/watch?v=yIYKR4sgzI8>

<https://www.analyticsvidhya.com/blog/2015/11/beginners-guide-on-logistic-regression-in-r/>

<https://www.dataschool.io/simple-guide-to-confusion-matrix-terminology/>

#### Conclusion:

- ❖ Logistic Regression is a classification model where we predict based on independent variables whether the classification is '0' or '1'(in binary)
- ❖ We have learnt about multi-variate(multiple independent variables) and single-variable regression and also multi-class logistic regression(multiple dependent classes)
- ❖ We have learnt about sigmoid function and realized how it helps for the hypothesis for classification of dependent variable
- ❖ We have learnt about various parameters which tell us about the accuracy of our predicted model using confusion matrix
- ❖ We have learnt how to build a model for our data,train it and then use it for prediction purposes.